

Introduction to Parallel Computing

a.y. 2025 – 2026

Dr. Brahmaiah Gandham

Lecture 10
Shared Memory programming with OpenMP

Outline

- **Data Sharing Clauses**
- **Work Sharing Directives**
- **Synchronization Directives**

Data Sharing Clauses

- **private**

- each thread gets a private copy of variable, no longer “storage-associated” with original variable
- private copy is not initialized
- default for variables declared inside parallel region
- often better to declare variables inside parallel region, reduces amount of code
(also minimizes “vertical distance” in source code)

```
int main() {
    int x = 10;

#pragma omp parallel private(x)
{
    int tid = omp_get_thread_num();
    x = 10; // must initialize, each thread has its own copy
    x = x + tid;
    printf("Thread %d: x = %d\n", tid, x);
}

printf("After parallel region: x = %d\n", x);
return 0;
}
```

Data Sharing Clauses

- **shared**

- each thread references the same, global copy
- data races if access is not synchronized
- default for variables declared outside parallel region and global variables, often used for read-only access

```
int main() {
    int x = 10;

    #pragma omp parallel shared(x)
    {
        int tid = omp_get_thread_num();
        x = x + tid; // all threads modify the same variable
        printf("Thread %d: x = %d\n", tid, x);
    }

    printf("After parallel region: x = %d\n", x);
    return 0;
}
```

Data Sharing Clauses

- **default**
 - can be set to shared, or none for C/C++
 - none helpful for detecting missing variables in clauses (compiler will complain!)

```
int main() {
    int x = 10;

    // Must explicitly specify every variable
    #pragma omp parallel default(none) private(x)
    {
        int tid = omp_get_thread_num();
        x = 10; // each thread has its own copy
        x = x + tid;
        printf("Thread %d: x = %d\n", tid, x);
    }

    printf("After parallel region: x = %d\n", x);
    return 0;
}
```

Data Sharing Clauses cont'd

- `firstprivate`
 - like `private`, but private copies are initialized with value of copy outside of parallel region

```
int x = 10;

#pragma omp parallel firstprivate(x)
{
    x += omp_get_thread_num(); // each thread starts with x=10
    printf("Thread %d: x = %d\n", omp_get_thread_num(), x);
}

printf("After parallel: x = %d\n", x); // still 10
```

Data Sharing Clauses cont'd

- `lastprivate`
 - like `private`, but outside copy is set to the private copy of the final iteration (for loops) or last section (`sections`), **NOT** the iteration/section that was chronologically executed last

```
int x = 0;

#pragma omp parallel for lastprivate(x)
for (int i = 0; i < 4; i++) {
    x = i * 10;
    printf("Thread %d: x = %d\n", omp_get_thread_num(), x);
}

printf("After parallel for: x = %d\n", x);
```

Data Sharing Clauses cont'd

- `threadprivate`
 - like `private`, but will persist across parallel regions
 - master thread variable is storage-associated with original variable (not the case for `private`!)

```
static int counter = 0;  
#pragma omp threadprivate(counter)  
  
#pragma omp parallel  
{  
    counter = omp_get_thread_num();  
    printf("Region 1 - Thread %d: counter = %d\n", omp_get_thread_num(), counter);  
}  
  
#pragma omp parallel  
{  
    printf("Region 2 - Thread %d: counter = %d\n", omp_get_thread_num(), counter);  
}
```

Reduction Clause

- performs reduction to a single variable in parallel or loop context
 - arithmetic ops: +, -, *, max, min
 - logical ops: &, &&, |, ||, ^
 - careful with associativity of floating-point operations!
- user-defined reductions are possible (version 4.0)
 - need to be declared with `#pragma omp declare reduction`

```
#pragma omp parallel
{
    #pragma omp for reduction(+:x)
    for(int i = 0; i < 10; ++i) {
        x += i;
    }
}

// or

#pragma omp parallel reduction(-:x)
x -= omp_get_thread_num();
```

Work Sharing Directives

- distribute execution of following code region among existing threads of the team
- must be enclosed in parallel region, cannot be directly nested
- do not launch new threads but assign work to existing threads
- no barrier on entry
- implicit barrier on exit
 - **unless nowait clause specified!**
- for
- sections
- single
- task
- simd

sections Directive

- sections may be executed concurrently, each by an arbitrary thread of the team
- matches MIMD programming patterns
- easily leads to load imbalance if individual sections not equally work-intensive
 - also, maximum degree of parallelism limited by number of sections

```
#pragma omp parallel
{
    #pragma omp sections
    {
        #pragma omp section
        { ... }
        #pragma omp section
        { ... }
        #pragma omp section
        { ... }
    }
}
```

single Directive

- code region will only be executed by a single, arbitrary thread
 - useful for interacting with libraries, that are not multi-threading-aware
- implicit barrier at the end for all threads in the team
- also available as master variant
 - like single, but for master thread
 - no implicit barrier at the end

```
#pragma omp parallel
{
    #pragma omp single
    {
        ...
    }
}
```

Synchronization Directives

- constructs discussed so far are fairly high-level in terms of synchronization
 - e.g. unable to enforce specific order or simply wait for all threads
- OpenMP also offers more fine-grained synchronization directives
 - barrier
 - critical
 - atomic
 - ordered
 - flush

barrier Directive

- explicit barrier requested by user
- threads are not allowed to continue until all have reached the barrier
- Implicit barrier at the end of for, sections, single, task, simd unless nowait specified

```
#pragma omp parallel
{
    ...
#pragma omp barrier
    ...
}
```

critical Directive

code region executed by all threads but only one at a time

can be named to allow finer-grained mutual exclusion

only critical regions with the same name enforce mutual exclusion

all regions without a name have the same name

```
#pragma omp parallel
{
    #pragma omp critical(foo)
    { ... }
    #pragma omp critical(foo)
    { ... }
    #pragma omp critical(bar)
    { ... }
    #pragma omp critical
    { ... }
    #pragma omp critical
    { ... }
}
```

atomic Directive

- same as `critical`, but restricted to a single memory location and certain operations
- restriction allows mapping to fast hardware mechanisms
 - optional specialization clauses `read`, `write`, `update`, and `capture`
- keeps code hardware- and compiler-independent compared to using intrinsics
 - but may just be a wrapper for `critical` e.g. when lacking hardware support

```
int count = 0;  
#pragma omp parallel  
{  
    #pragma omp atomic  
    count++;  
}
```

ordered Directive + Clause

- enforces critical region with sequential execution order
- required for strong sequential equivalence
 - but at high performance cost
- only efficient if code outside the ordered region is expensive enough
 - remember guidelines about strong / weak sequential equivalence
 - check whether your numerical method really needs strong sequential equivalence

```
double total = 0.0;  
#pragma omp parallel for ordered  
for(int i = 0; i < N; ++i) {  
    double part = f(i);  
    #pragma omp ordered  
    total += part;  
}
```

flush Directive

problem: writes by one thread are not immediately visible to another in the same parallel region without synchronization

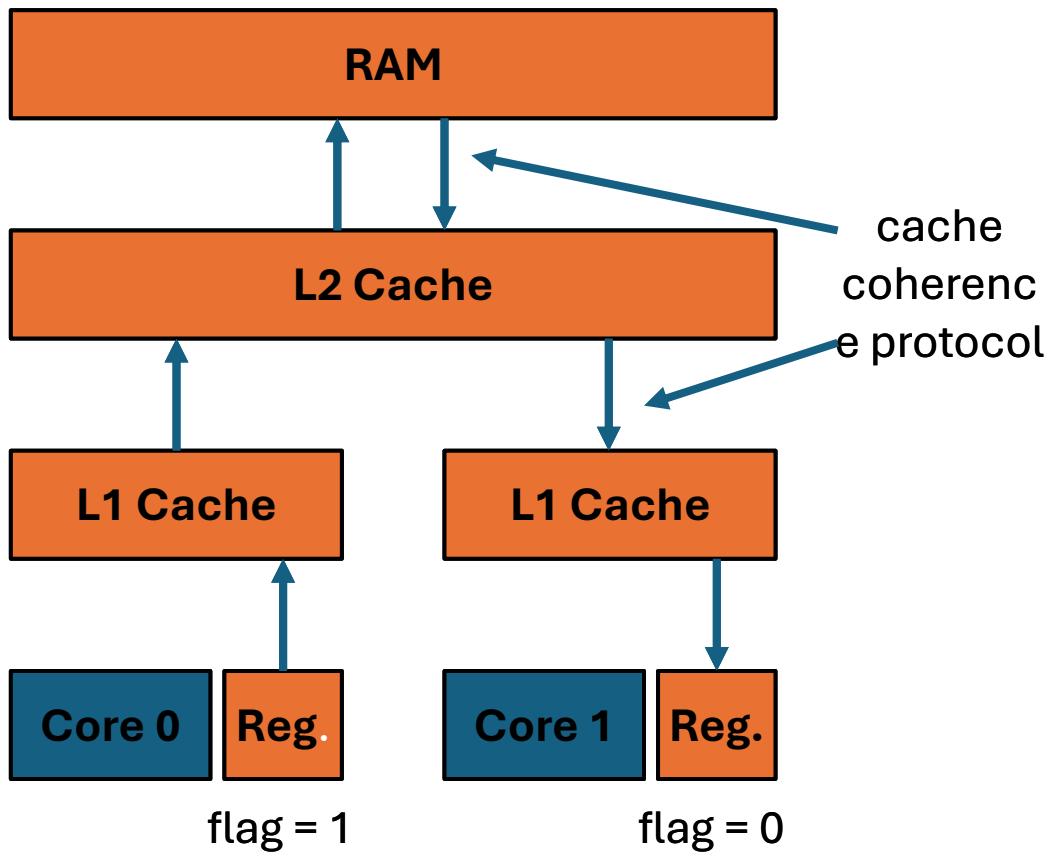
Example:

the second section could hang in while the loop

- But why?

```
#pragma omp parallel sections
{
    #pragma omp section
    { // producer section
        // produce some data
        flag = 1;
    }
    #pragma omp section
    { // consumer section
        while (flag == 0) { }
        // use data
    }
}
```

Detour: Register Spilling (or Lack Thereof)



```
#pragma omp parallel sections
{
    #pragma omp section
    { // producer section
        // produce some data
        flag = 1;
    }
    #pragma omp section
    { // consumer section
        while (flag == 0) { }
        // use data
    }
}
```

flush Directive cont'd

```
#pragma omp parallel sections
{
    #pragma omp section
    { // producer section
        // produce and flush data
        #pragma omp flush
        #pragma omp atomic write
        flag = 1;
        #pragma omp flush(flag)
    }
}
```

```
#pragma omp section
{ // consumer section
    while (1) {
        #pragma omp flush(flag)
        #pragma omp atomic read
        int temp_flag = flag;
        if(temp_flag == 1) break;
    }
    // use data
}
} // end sections
```

flush Directive cont'd

- flush provides a consistent view of the data across threads
- it is implied at
 - barrier
 - entry and exit of **critical**
 - exit of **parallel**, **for**, **sections**, **single**
 - set/unset of locks
- if otherwise required, use flush directive explicitly but with care (performance impact!)

Summary

- main characteristics
 - incremental parallelization
- programming, execution and memory models
 - based on threads and shared data access
 - mainly relies on pragmas as programmer interface
- directives
 - parallelism, work sharing, data sharing, synchronization
- reference material
 - “Parallel Programming for Science and Engineering” by Victor Eijkhout,
<https://web.corral.tacc.utexas.edu/CompEdu/pdf/pcse/EijkhoutParallelProgramming.pdf>